



Waqas Rasheed Khan

2295250

Logbook

Principles of Data Mining and Machine Learning  
School of Computing and Information Sciences

# LAB 1

The pandas library, in Python, provides a `describe()` function that gives an overview of the statistics for a `DataFrame`. It offers details like count, standard deviation, minimum and maximum values, for data. This functionality allows us to analyze and understand the distribution of data.

```
In [8]: df.describe()
```

```
Out[8]:
```

	Account length	Area code	Number vmail messages	Total day minutes	Total day calls	Total day charge	Total eve minutes	Total eve calls	Total eve charge	Total night minutes	Total night calls	Total night charge	Total intl minutes	Total intl calls	Total intl charge	Customer service calls
count	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00
mean	101.06	437.18	8.10	179.78	100.44	30.56	200.98	100.11	17.08	200.87	100.11	9.04	10.24	4.48	2.76	1.56
std	39.82	42.37	13.69	54.47	20.07	9.26	50.71	19.92	4.31	50.57	19.57	2.28	2.79	2.46	0.75	1.32
min	1.00	408.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	23.20	33.00	1.04	0.00	0.00	0.00	0.00
25%	74.00	408.00	0.00	143.70	87.00	24.43	166.60	87.00	14.16	167.00	87.00	7.52	8.50	3.00	2.30	1.00
50%	101.00	415.00	0.00	179.40	101.00	30.50	201.40	100.00	17.12	201.20	100.00	9.05	10.30	4.00	2.78	1.00
75%	127.00	510.00	20.00	216.40	114.00	36.79	235.30	114.00	20.00	235.30	113.00	10.59	12.10	6.00	3.27	2.00
max	243.00	510.00	51.00	350.80	165.00	59.64	363.70	170.00	30.91	395.00	175.00	17.77	20.00	20.00	5.40	9.00

## Read file

```
df = pd.read_csv("telecom_churn.csv")
df.head()
```

```
In [2]: df=pd.read_csv("telecom_churn.csv")
df.head()
```

```
Out[2]:
```

	State	Account length	Area code	International plan	Voice mail plan	Number vmail messages	Total day minutes	Total day calls	Total day charge	Total eve minutes	Total eve calls	Total eve charge	Total night minutes	Total night calls	Total night charge	Total intl minutes	Total intl calls	Total intl charge	Customer service calls
0	KS	128	415	No	Yes	25	265.1	110	45.07	197.4	99	16.78	244.7	91	11.01	10.0	3	2.70	
1	OH	107	415	No	Yes	26	161.6	123	27.47	195.5	103	16.62	254.4	103	11.45	13.7	3	3.70	
2	NJ	137	415	No	No	0	243.4	114	41.38	121.2	110	10.30	162.6	104	7.32	12.2	5	3.29	
3	OH	84	408	Yes	No	0	299.4	71	50.90	61.9	88	5.26	196.9	89	8.86	6.6	7	1.78	
4	OK	75	415	Yes	No	0	166.7	113	28.34	148.3	122	12.61	186.9	121	8.41	10.1	3	2.73	

```
In [8]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   State                                3333 non-null   object
1   Account length                       3333 non-null   int64
2   Area code                           3333 non-null   int64
3   International plan                   3333 non-null   object
4   Voice mail plan                     3333 non-null   object
5   Number vmail messages               3333 non-null   int64
6   Total day minutes                   3333 non-null   float64
7   Total day calls                     3333 non-null   int64
8   Total day charge                    3333 non-null   float64
9   Total eve minutes                   3333 non-null   float64
10  Total eve calls                     3333 non-null   int64
11  Total eve charge                    3333 non-null   float64
12  Total night minutes                 3333 non-null   float64
13  Total night calls                   3333 non-null   int64
14  Total night charge                  3333 non-null   float64
15  Total intl minutes                  3333 non-null   float64
16  Total intl calls                    3333 non-null   int64
17  Total intl charge                   3333 non-null   float64
18  Customer service calls              3333 non-null   int64
19  Churn                              3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(3)
memory usage: 498.1+ KB
None
```

You have the ability to organize a dataset according to the value of one of its variables, such as columns. For example you can sort it by day charge by utilizing `ascending=False` to arrange it in descending order.

```
In [17]: df.sort_values(by="Total day charge", ascending = False).head()
```

Out[17]:

	State	Account length	Area code	International plan	Voice mail plan	Number vmail messages	Total day minutes	Total day calls	Total day charge	Total eve minutes	Total eve calls	Total eve charge	Total night minutes	Total night calls	Total night charge	Total intl minutes	Total intl calls	Total intl charge	C
365	CO	154	415	No	No	0	350.8	75	59.64	216.5	94	18.40	253.9	100	11.43	10.1	9	2.73	
985	NY	64	415	Yes	No	0	346.8	55	58.96	249.5	79	21.21	275.4	102	12.39	13.3	9	3.59	
2594	OH	115	510	Yes	No	0	345.3	81	58.70	203.4	106	17.29	217.5	107	9.79	11.8	8	3.19	
156	OH	83	415	No	No	0	337.4	120	57.36	227.4	116	19.33	153.9	114	6.93	15.8	7	4.27	
605	MO	112	415	No	No	0	335.5	77	57.04	212.5	109	18.06	265.0	132	11.93	12.7	8	3.43	

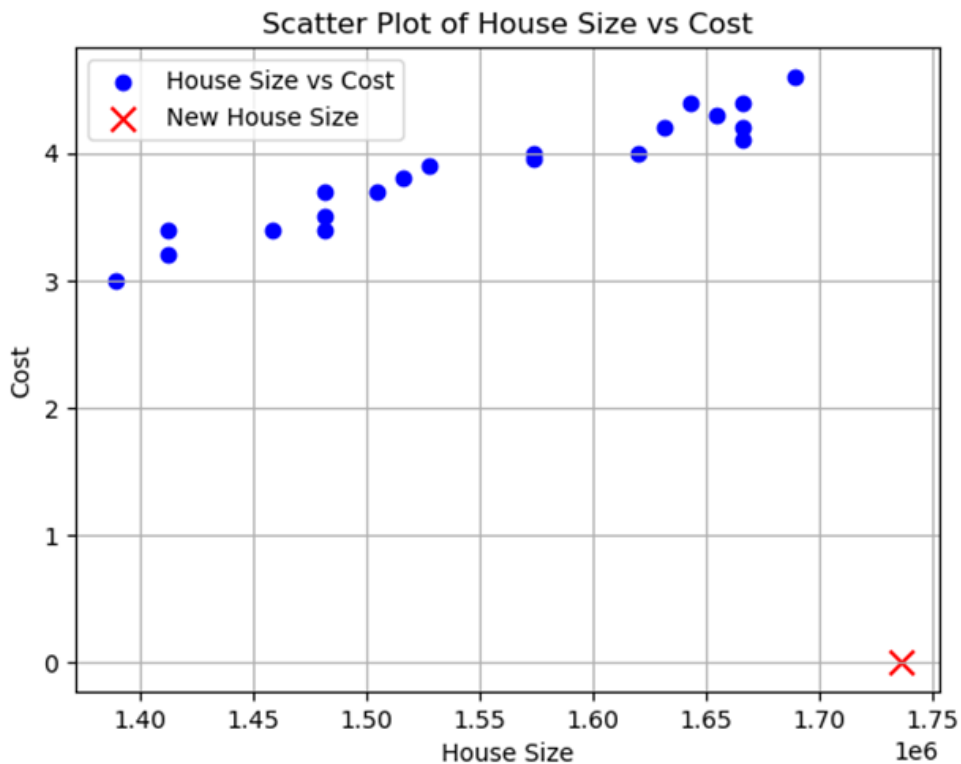
In Python the pandas library offers a function called `value_counts()`. This function allows you to determine how times each distinct value occurs in a Series.

```
In [14]: df["International plan"].value_counts()
```

```
Out[14]: No      3010  
        Yes       323  
        Name: International plan, dtype: int64
```

## Lab 2


A Scatter plot is commonly used to show the correlation between the size of a house and its cost. The Y Axis represents the cost while the X Axis indicates the size of the house. It can be observed that as the size of a house increases its cost also tends to increase.



## Lab 3

```
In [22]: # Your code to plot all costs
plt.plot(J_history)
plt.title('Iteration vs Cost - Line Plot')
```

Out[22]: Text(0.5, 1.0, 'Iteration vs Cost - Line Plot')



```
In [37]: SID = 2295250
First_City = SID/10 # Put the population of first city as 10 times Less than your SID
Second_City = SID/30 # Put the population of second city as 30 times Less than your SID

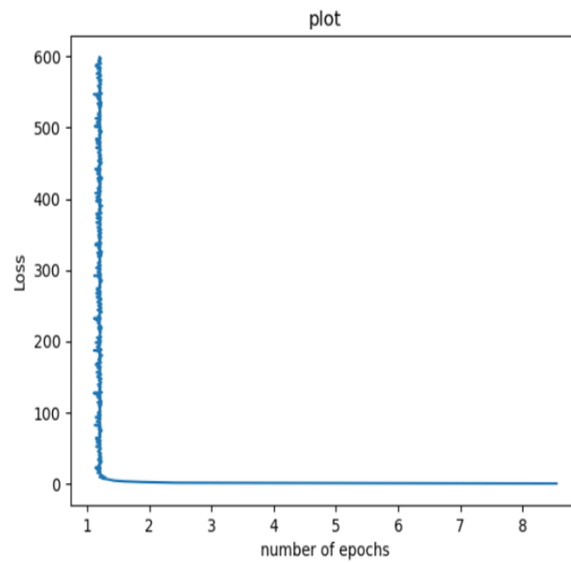
predict1 = (prediction([1, First_City/10000]),(new_theta))
predict2 = (prediction([1, Second_City/10000]),(new_theta))

print(f'For a population of {First_City} people, profit will be {predict1[0]} ')
print(f'For a population of {Second_City} people, profit will be {predict2[0]}')
```

For a population of 229525.0 people, profit will be 234641.7190432843  
For a population of 76508.33333333333 people, profit will be 52360.23184742697

## Lab 4

```
In [88]: # Your code to plot epochs vs loss. Call the method.  
plotLoss(num_epochs,train_loss)
```



```
In [89]: #Code to predict the profit i.e. y values  
predict(theta_updated,y_train)
```

```
Out[89]: -4.797723428027043
```

## Lab 5

**Calculating accuracy - number of correct classification/total number of classification.**

```
In [42]: np.mean(predict(res.x, X) == y)
```

```
Out[42]: 0.89
```

```
In [22]: resultant_accuracy = 0.89
SID = 2295250
encrypted_value = resultant_accuracy*SID
print(encrypted_value)


2042772.5
```

## Lab 6

In my analysis I have evaluated the precision of MLP models using the class. I considered three accuracy levels, 40%, 60% and 25%.

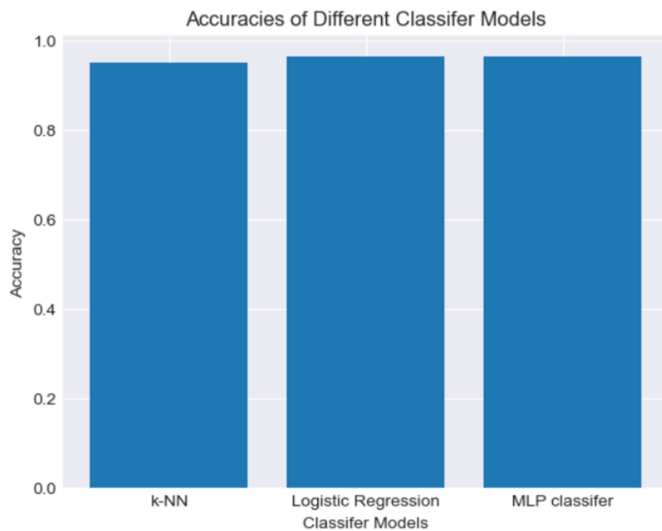
```
In [17]: #Assuming accuracy - 40, 60, 25
plt.barh(['acc_1', 'acc_2', 'acc_3'], [40, 60, 25])

Out[17]: <BarContainer object of 3 artists>
```

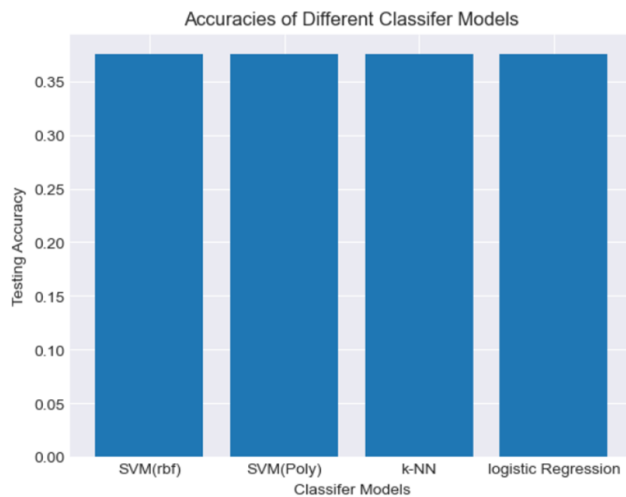


## Lab 8

In this case a bar chart is employed to showcase the variation, in accuracies across classifier models such, as KNN, Logistic Regression and MLP Classifier. It appears that there is no disparity observed among these classifier models.



## Lab 9



According to the confusion matrix logistic regression demonstrates the performance while KNN performs the least effectively.



# Lab 10

Q1:

At each node of the decision tree the criteria used to make decisions are information gain and Gini impurity. These criteria help determine the quality of test decisions and how they classify samples into classes.

Q2:

In a decision tree the measure of disorder such as entropy, Gini index or loss decreases at each node when the split leads to subsets that are more similar. The objective is to reduce disorder (entropy or Gini index) or loss in order to improve the tree's ability to differentiate or make predictions about the target variable.

Q3:

As we descend further into the tree it is common for the values of Entropy/Gini/Loss Change to decrease.


Q4:

Typically the quantity of data samples at each node in a tree building process is influenced by the data itself and the splits that are made. A general trend is that nodes representing categories or conditions which are closer to the root tend to have a higher number of samples.

Q5:

In a decision tree when we reach a leaf node it usually contains information about the predicted output or class for the subset of data that led to that leaf. In classification tasks the leaf node holds the class label assigned to the majority of samples, within that subset.

# Lab 11



```
In [38]: kmeans = KMeans(n_clusters=3, random_state=0)
kmeans.fit(X)
# check how many of the samples were correctly labeled
labels = kmeans.labels_
correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))
```

C:\Users\asimq\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1412: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning  
super().\_check\_params\_vs\_input(X, default\_n\_init=10)

Result: 4165 out of 7050 samples were correctly labeled.  
Accuracy score: 0.59

```
In [ ]: kmeans = KMeans(n_clusters=4, random_state=0)
kmeans.fit(X)
# check how many of the samples were correctly labeled
labels = kmeans.labels_
correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))
```

Result: 4340 out of 7050 samples were correctly labeled.  
Accuracy score: 0.62

